

Big Data Foundations

- Sai Kiran Krishna Murthy

Agenda

- Day 1: Introduction
- Day 2: HDFS, Kafka, Spark
- Day 3: Practical Showcase
- Day 4: Free day

What is Big Data?

The 5 V's of Big Data

- Volume
- Velocity
- Variety
- Veracity
- Value

Big Data is a Paradigm shift

- Schema on write vs Schema on read
- Scale up vs Scale out
 - Horizontal Scaling vs Vertical Scaling
 - Single Node vs Cluster Computing
- Specialized vs Commodity hardware

Schema on write

- Create table
- Write data (validation happens here)
- Read data

Schema on read

- Create file (Write data)
- Read data (validation happens here)

Scale up vs Scale out

Scale up

- Get a bigger machine
 - More RAM
 - Bigger CPU/GPU
 - More Cores



Problems?

- Increasing costs
- Physical Limitation
- Specialized software? → Vendor Lockin

Alternate?

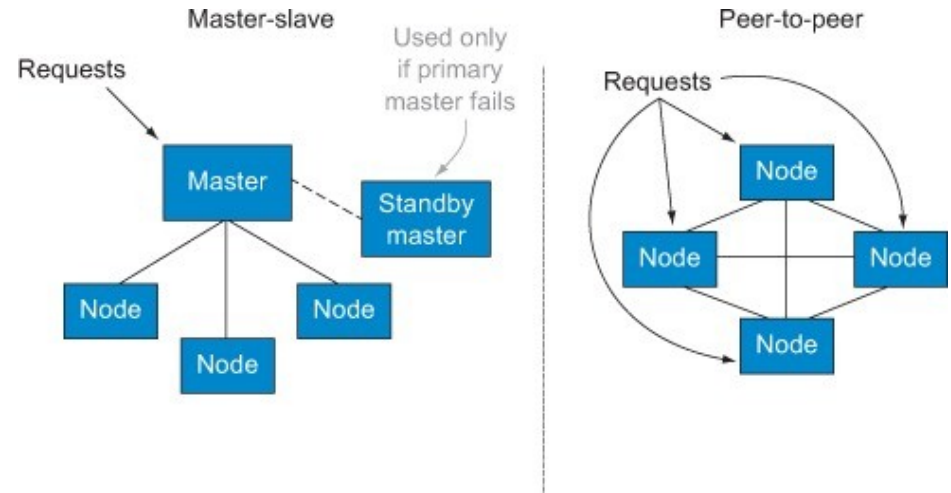
- Scale out:
 - add more machines (nodes)
 - Create a cluster of these machines (nodes)

Scale out: In practice



Features

- Concurrency
- Scalability
- Availability
- Fault Tolerance



Case Study



Problem Statement

- Most of Scrooge's wealth is in his Gold Coins
- Total wealth \sim Number of Gold Coins * Cost of each gold coin



Introducing: The Counter

- He counts gold coins for a living



Current Approach

- Identify gold coins from the vault
- Move the gold to where the counter is
- Let the counter do the counting and give us the results
- Currently we have 1 counter

- Total time taken
 - Sorting and Filter time
 - Transportation time
 - Counting time

How can we Improve?

Aspects we can improve?

- Storage Layout: How we store?
- Locality: Where we count?
- Processing Paradigm: How we count?

Storage Layout

Scenario #1: Row based



Discussion #1

- Advantages and disadvantages

Storage Layout

Scenario #2: Column based



Discussion #2

- Advantages and disadvantages

Storage Layout

Scenario #3: Hybrid



Locality

- Move Gold to where the Counter is?
- Move the Counter to the Vault?

How we count? (Processing Paradigm)

- Divide and Conquer
 - Divide the data
 - Add more counters
 - Let the counters work in parallel

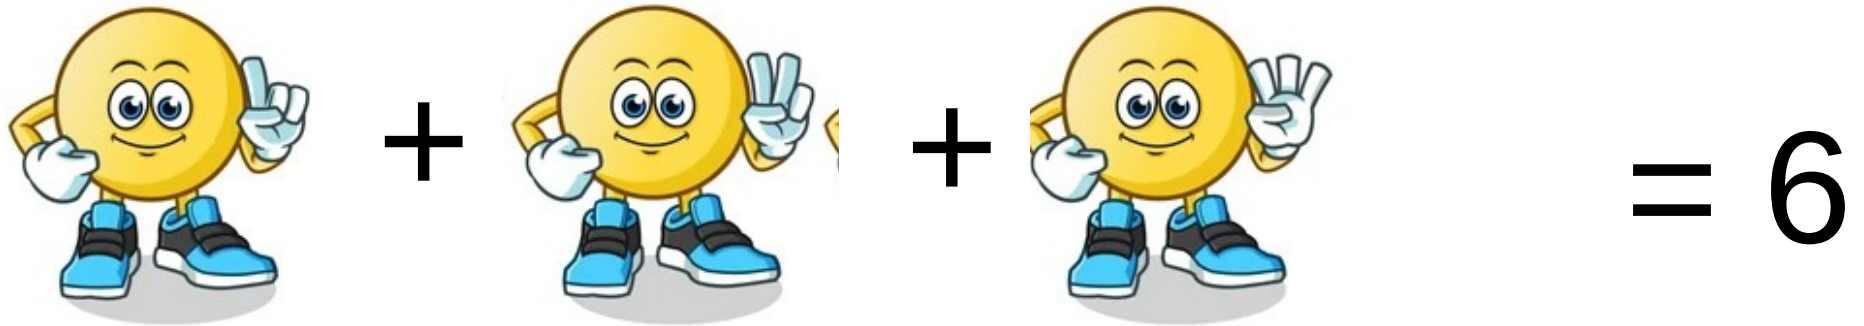
Bucketing / Partitioning



Mapping



Reducing



A visual equation illustrating the concept of reducing. It features three identical yellow cartoon characters, each with a single finger held up. These characters are separated by plus signs. To the right of the third character is an equals sign followed by the number 6. This visualizes the process of reducing three individual units into a single, simplified unit of 6.

$$1 + 1 + 1 = 6$$

Map + Reduce

- Split a big problem in smaller pieces
- Map each task to a counter (worker)
- Reduce the results of each worker, until you have only one result

Discussion

- What happens with too few partitions?
- What happens with too many partitions?
- What is the ideal number of partitions?

Discussion

- What happens with too few counters?
- What happens with too many counters?
- What is the ideal number of counters?

Examples

Exercise

- Take 10 mins and come up with an example problem you can solve with map reduce
- Present it to the group

Wait a minute

- Were we talking about Big Data all along?
 - Gold == Data
 - Scrooge has a lot of data (aka volume) == Big Data
 - Counter == A program
 - Vault == Storage
 - MapReduce == processing framework
 - Different Binary File formats == avro, parquet

Recap Lessons learned

- Data Locality : Move the program to where the data resides
- Program Paradigm: Map Reduce
- Data Storage: Depending on the use case, Row Vs Column Storage

Introducing Hadoop

- Hadoop = Storage + Processing Framework + Scheduler
 - Storage = HDFS
 - Processing Framework = Map Reduce
 - Scheduler = Yarn

Note on Map Reduce

- Map Reduce has been replaced with Apache Spark

Putting it all together

- Storage = HDFS
- Processing Framework = Spark
- Scheduler = Yarn
- Synchronization = Zoo Keeper
- SQL interface for HDFS = Hive
- NoSQL Database = HBase
- Message Queue = Kafka
- Notebook = Zeppelin

CDP: Walk through

Extra slides

Avro file format

- Rich data types
- Row based format
- Binary
- Splittable

Introducing CAP theorem

C – Consistency

A – Availability

P – Partition Tolerance

Distributed System?
You can only have
two

Consistency

C

All Clients have same view
of Data

- . Hbase
- . Google Big Table
- . Mongo DB
- . Redis
- . HyperTable

- . RDBMS(Oracle,
MySQL)
- . Vertica
- . Aster Data

Pick Two

System is Functional in
spite of network Partition

P

Partition
Tolerance

- . Dynamo DB
- . Cassandra
- . Voldemort

- . Riak
- . Couch DB

A

Availability

All Clients Can Read and
Write